

No-Bullshit Data Science

Szilárd Pafka, PhD

Chief Scientist, Epoch

R/Finance Conference

Chicago, May 2017



Szilard

@ DataScienceLA
Santa Monica, California

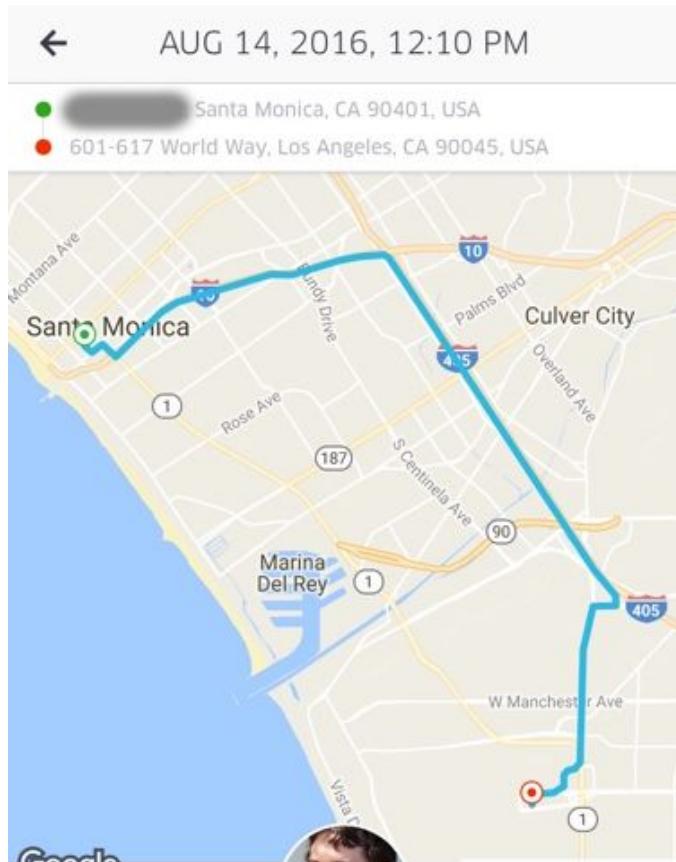
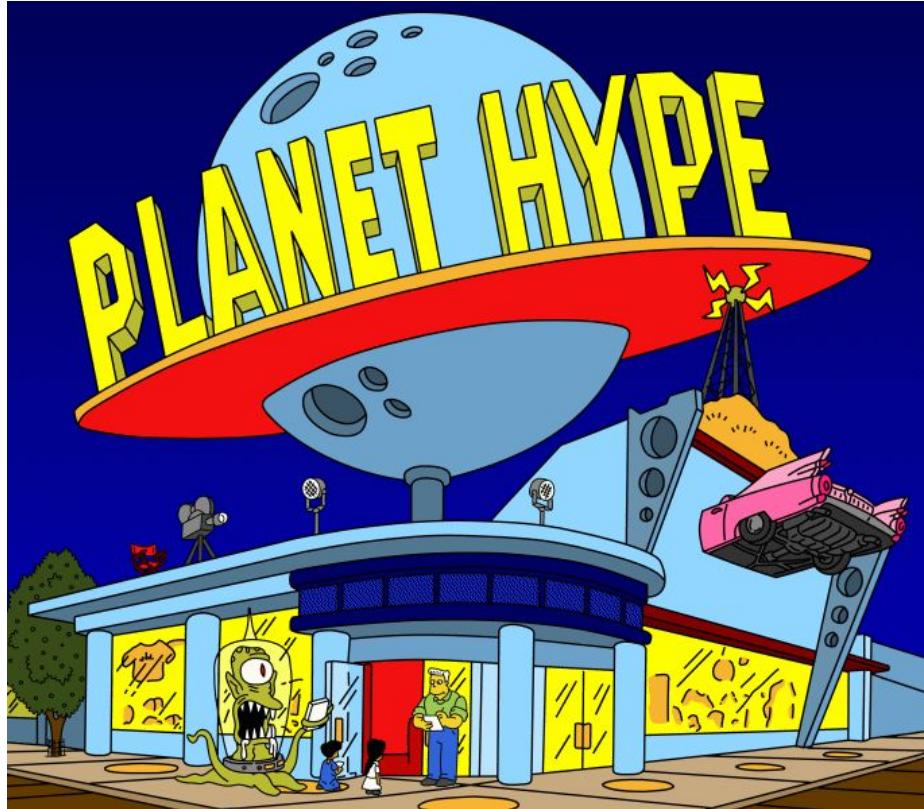
physics PhD, chief (data) scientist, meetup
organizer, datascience.la, visiting professor

Disclaimer:

I am not representing my employer (Epoch) in this talk

I cannot confirm nor deny if Epoch is using any of the methods, tools, results etc. mentioned in this talk





YOU RATED
★★★★★



COLIN

TRIP TOTAL
\$10.61

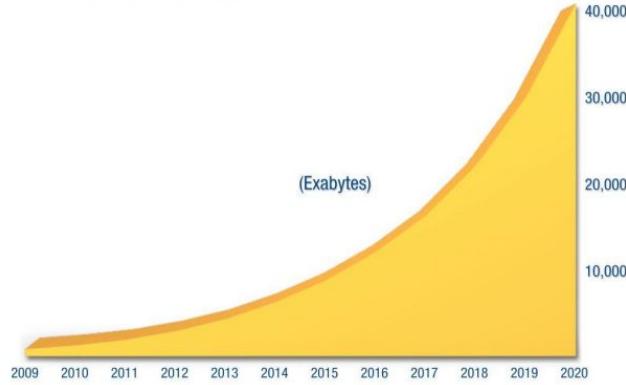




Example #1

Figure 1

The Digital Universe: 50-fold Growth from the Beginning of
2010 to the End of 2020

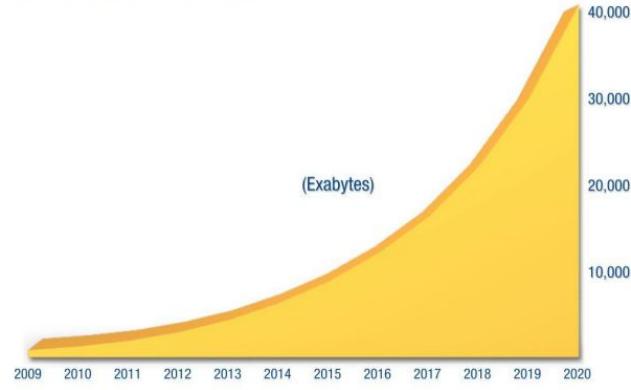


Source: IDC's Digital Universe Study, sponsored by EMC, December 2012



Figure 1

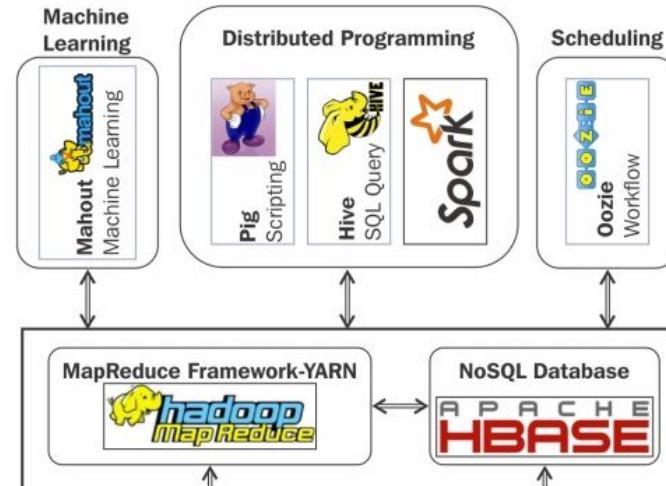
The Digital Universe: 50-fold Growth from the Beginning of 2010 to the End of 2020



Source: IDC's Digital Universe Study, sponsored by EMC, December 2012



imgflip.com







Gary Bernhardt
@garybernhardt



Consulting service: you bring your big data problems to me, I say "your data set fits in RAM", you pay me \$10,000 for saving you \$500,000.

RETWEETS
997

LIKES
960



3:03 PM - 19 May 2015



Kingston Technology Value RAM 128GB Kit (4x32GB) 2133MHz DDR4 ECC Reg CL15 (KVR21R15D4K4/128)

by [Kingston Technology](#)

[Be the first to review this item](#)

Was: \$743.99

Price: **\$743.96** & FREE Shipping. [Details](#)





Kingston Technology Value RAM 128GB Kit (4x32GB) 2133MHz DDR4 ECC Reg CL15 (KVR21R15D4K4/128)

by [Kingston Technology](#)

[Be the first to review this item](#)

Was: \$743.99

Price: **\$743.96** & FREE Shipping. [Details](#)



Model	vCPU	Mem (GiB)
-------	------	-----------

r3.8xlarge	32	244
------------	----	-----

x1.32xlarge	128	1,952
-------------	-----	-------

x1.32xlarge

128

349

1952

2 x 1920 SSD

\$13.338 per Hour

1	[██████████]	89.6%	33	[██████████]	88.5%	65	[██████████]	82.8%	97	[██████████]	88.5%
2	[██████████]	85.8%	34	[██████████]	88.5%	66	[██████████]	85.3%	98	[██████████]	89.7%
3	[██████████]	91.1%	35	[██████████]	84.0%	67	[██████████]	86.5%	99	[██████████]	86.5%
4	[██████████]	85.3%	36	[██████████]	85.9%	68	[██████████]	85.4%	100	[██████████]	88.5%
5	[██████████]	84.6%	37	[██████████]	89.2%	69	[██████████]	85.4%	101	[██████████]	84.7%
6	[██████████]	84.6%	38	[██████████]	89.1%	70	[██████████]	86.5%	102	[██████████]	88.5%
7	[██████████]	84.1%	39	[██████████]	87.3%	71	[██████████]	86.5%	103	[██████████]	89.2%
8	[██████████]	85.2%	40	[██████████]	87.9%	72	[██████████]	87.9%	104	[██████████]	86.5%
9	[██████████]	84.7%	41	[██████████]	83.4%	73	[██████████]	89.1%	105	[██████████]	87.1%
10	[██████████]	85.4%	42	[██████████]	85.4%	74	[██████████]	84.7%	106	[██████████]	84.6%
11	[██████████]	85.3%	43	[██████████]	85.8%	75	[██████████]	86.0%	107	[██████████]	84.6%
12	[██████████]	86.0%	44	[██████████]	91.7%	76	[██████████]	86.5%	108	[██████████]	90.4%
13	[██████████]	85.9%	45	[██████████]	89.2%	77	[██████████]	85.4%	109	[██████████]	90.4%
14	[██████████]	83.3%	46	[██████████]	87.2%	78	[██████████]	86.0%	110	[██████████]	87.8%
15	[██████████]	84.6%	47	[██████████]	87.1%	79	[██████████]	86.6%	111	[██████████]	85.9%
16	[██████████]	85.4%	48	[██████████]	87.1%	80	[██████████]	86.7%	112	[██████████]	87.8%
17	[██████████]	83.3%	49	[██████████]	89.7%	81	[██████████]	85.8%	113	[██████████]	88.5%
18	[██████████]	85.4%	50	[██████████]	94.3%	82	[██████████]	86.6%	114	[██████████]	91.1%
19	[██████████]	82.7%	51	[██████████]	91.7%	83	[██████████]	89.1%	115	[██████████]	91.1%
20	[██████████]	81.5%	52	[██████████]	96.2%	84	[██████████]	85.9%	116	[██████████]	90.4%
21	[██████████]	81.9%	53	[██████████]	91.0%	85	[██████████]	85.3%	117	[██████████]	89.2%
22	[██████████]	84.1%	54	[██████████]	90.4%	86	[██████████]	85.9%	118	[██████████]	92.3%
23	[██████████]	83.2%	55	[██████████]	90.5%	87	[██████████]	87.2%	119	[██████████]	93.6%
24	[██████████]	85.4%	56	[██████████]	89.1%	88	[██████████]	85.9%	120	[██████████]	91.7%
25	[██████████]	86.0%	57	[██████████]	88.5%	89	[██████████]	88.5%	121	[██████████]	91.1%
26	[██████████]	87.3%	58	[██████████]	91.7%	90	[██████████]	84.2%	122	[██████████]	91.1%
27	[██████████]	83.3%	59	[██████████]	96.2%	91	[██████████]	84.8%	123	[██████████]	89.8%
28	[██████████]	84.0%	60	[██████████]	91.7%	92	[██████████]	84.6%	124	[██████████]	90.4%
29	[██████████]	87.3%	61	[██████████]	88.5%	93	[██████████]	86.5%	125	[██████████]	93.6%
30	[██████████]	84.6%	62	[██████████]	90.4%	94	[██████████]	87.3%	126	[██████████]	92.9%
31	[██████████]	86.5%	63	[██████████]	92.3%	95	[██████████]	86.6%	127	[██████████]	90.4%
32	[██████████]	86.5%	64	[██████████]	94.2%	96	[██████████]	86.5%	128	[██████████]	91.0%
Mem	[██████████]	1062263/1967522MB	Swp	[██████████]	0/0MB	Tasks: 43, 190 thr, 1075 kthr; 129 running					
						Load average: 58.14 21.01 11.40					
						Uptime: 18:32:29					

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
34941	ubuntu	20	0	1025G	1024G	1380	S	11195	53.3	1h54:07	./pmbw -f ScanRead64PtrSimpleLoop -p 128 -P 128
34999	ubuntu	20	0	1025G	1024G	1380	R	95.3	53.3	0:50.06	./pmbw -f ScanRead64PtrSimpleLoop -p 128 -P 128
35066	ubuntu	20	0	1025G	1024G	1380	R	93.4	53.3	0:50.02	./pmbw -f ScanRead64PtrSimpleLoop -p 128 -P 128
34963	ubuntu	20	0	1025G	1024G	1380	R	92.7	53.3	0:51.13	./pmbw -f ScanRead64PtrSimpleLoop -p 128 -P 128



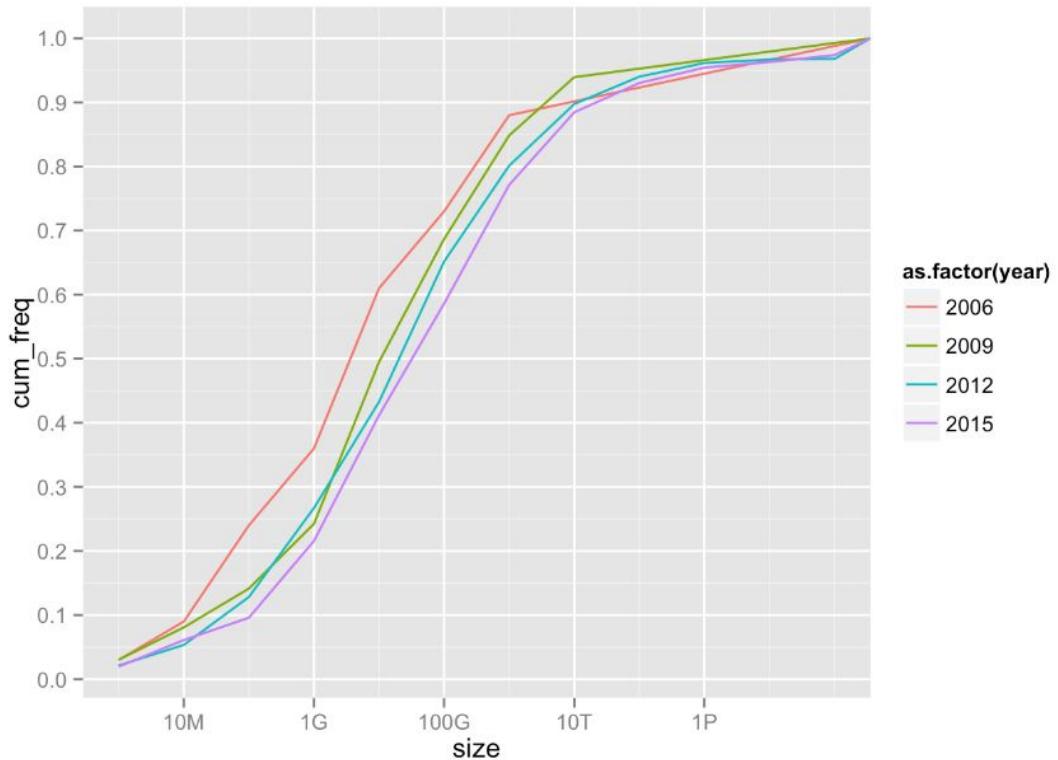
Largest Dataset Analyzed / Data Mined

[Tweet](#)

What was the largest dataset you analyzed / data mined? [392 votes]

2014 votes 2013 votes

szilard / dataset-sizes-kdnuggets





Szilard @DataScienceLA · Jun 17

What's the typical size of datasets you are analyzing?

18% <100MB

48% 100MB-10GB

18% 10GB-1TB

16% >1TB

151 votes • Final results



Szilard @DataScienceLA · Jun 17

What's the typical size of datasets you are analyzing?

18% <100MB

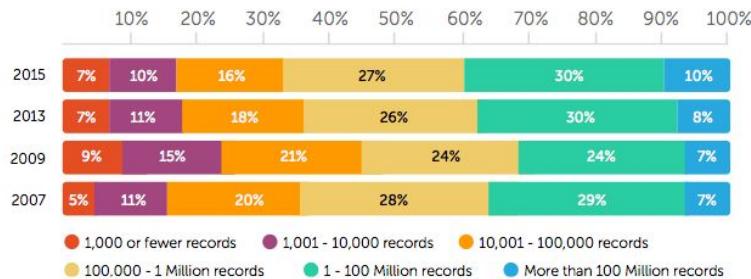
48% 100MB-10GB

18% 10GB-1TB

16% >1TB

151 votes • Final results

TYPICAL SIZE OF DATASETS





Szilard @DataScienceLA · Jun 17

What's the typical size of datasets you are analyzing?

18% <100MB

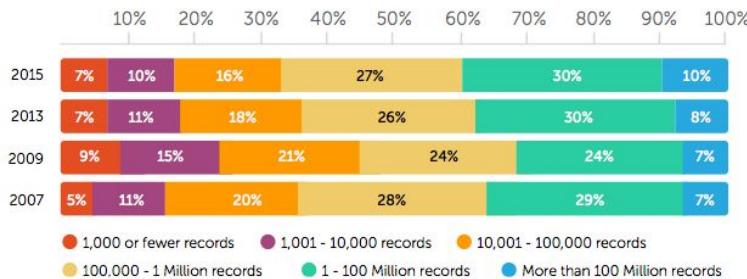
48% 100MB-10GB

18% 10GB-1TB

16% >1TB

151 votes • Final results

TYPICAL SIZE OF DATASETS



17 Rexer Analytics



Hadley Wickham

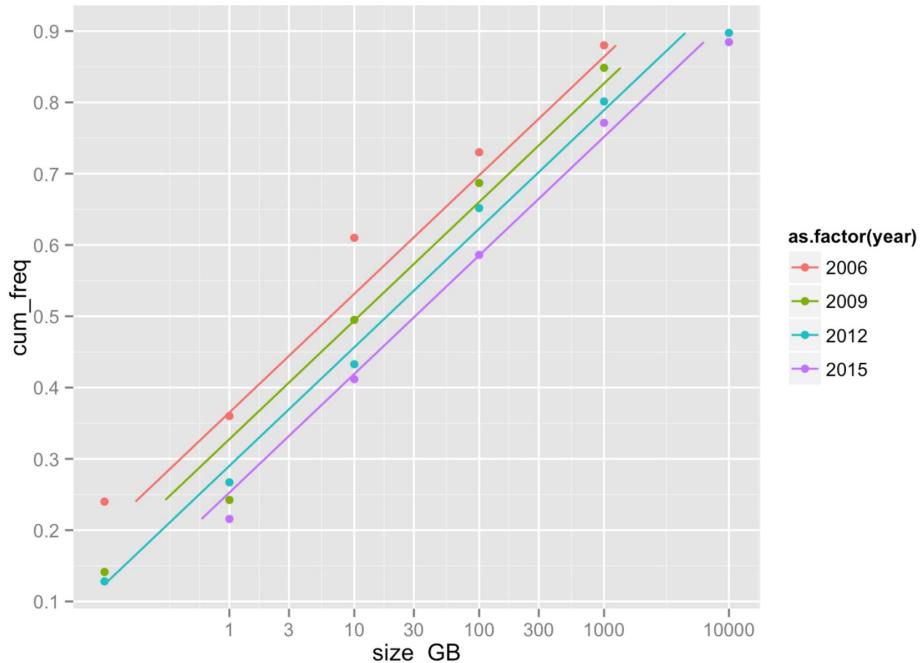
@hadleywickham



Following

"It takes a big man to admit his data is small" —
@jcheng

szilard / dataset-sizes-kdnuggets



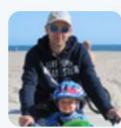
annual rate of increase of datasets of $10^{0.075} \sim 1.2$ that is 20%.

The [size of EC2 instances](#) with largest RAM:

year	type	RAM (GB)
2007	m1.xlarge	15
2009	m2.4xlarge	68
2012	hs1.8xlarge	117
2014	r3.8xlarge	244
2016*	x1	2 TB

The size of EC2 instances with largest RAM:

year	type	RAM (GB)
2007	m1.xlarge	15
2009	m2.4xlarge	68
2012	hs1.8xlarge	117
2014	r3.8xlarge	244
2016*	x1	2 TB



Szilard @DataScienceLA · 18 Nov 2015

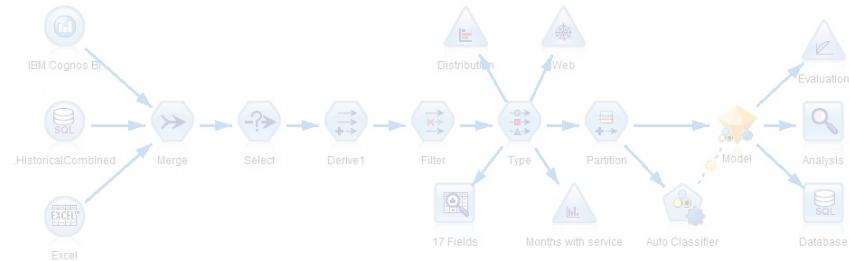
Big RAM is eating [#bigdata](#): datasets for analytics grew 20% /yr (last decade [@kdnuggets](#)), RAM EC2 grew 50% /yr



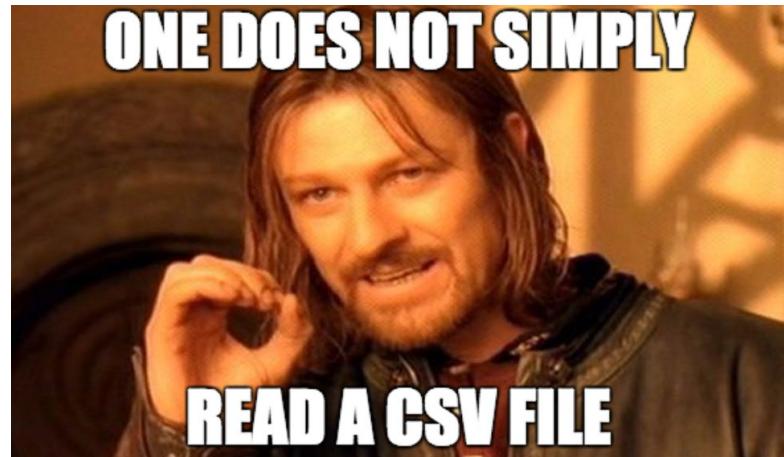




sas



data frame. Yesterday I wrote a [blog post](#) about using `sqldf()` to import the data into SQLite as a staging area, and then sucking it from SQLite into R. This works really well for me. I was able to pull in 2GB (3 columns, 40mm rows) of data in < 5 minutes. By contrast, the `read.csv` command ran all night and never completed.



answered Nov 30 '09 at 15:48



JD Long

27.5k • 31 • 135 • 216

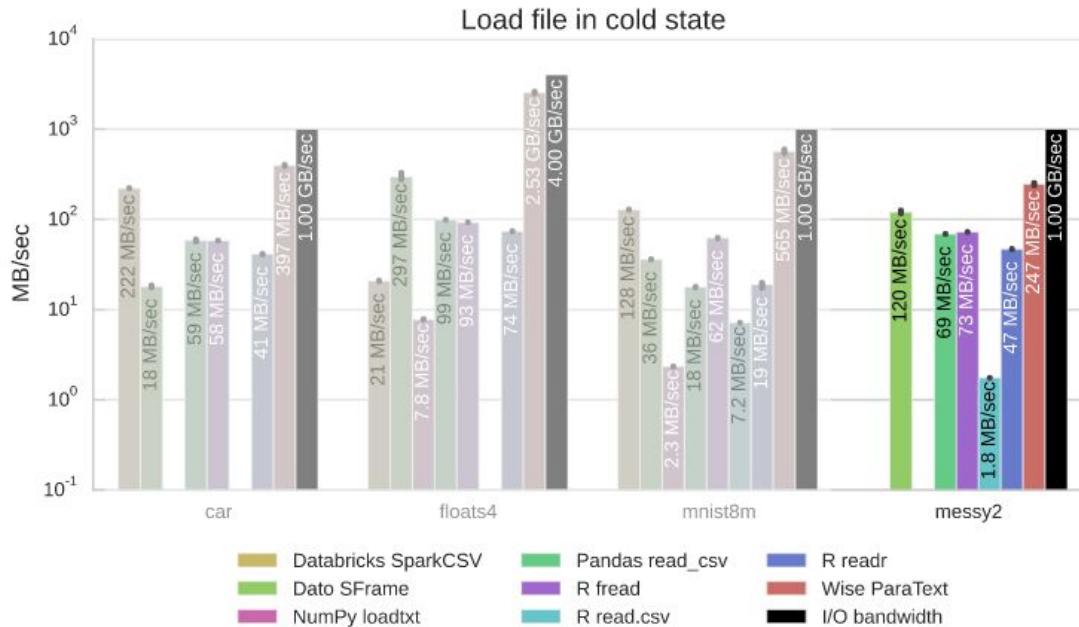
data frame. Yesterday I wrote a [blog post](#) about using `sqldf()` to import the data into SQLite as a staging area, and then sucking it from SQLite into R. This works really well for me. I was able to pull in 2GB (3 columns, 40mm rows) of data in < 5 minutes. By contrast, the `read.csv` command ran all night and never completed.

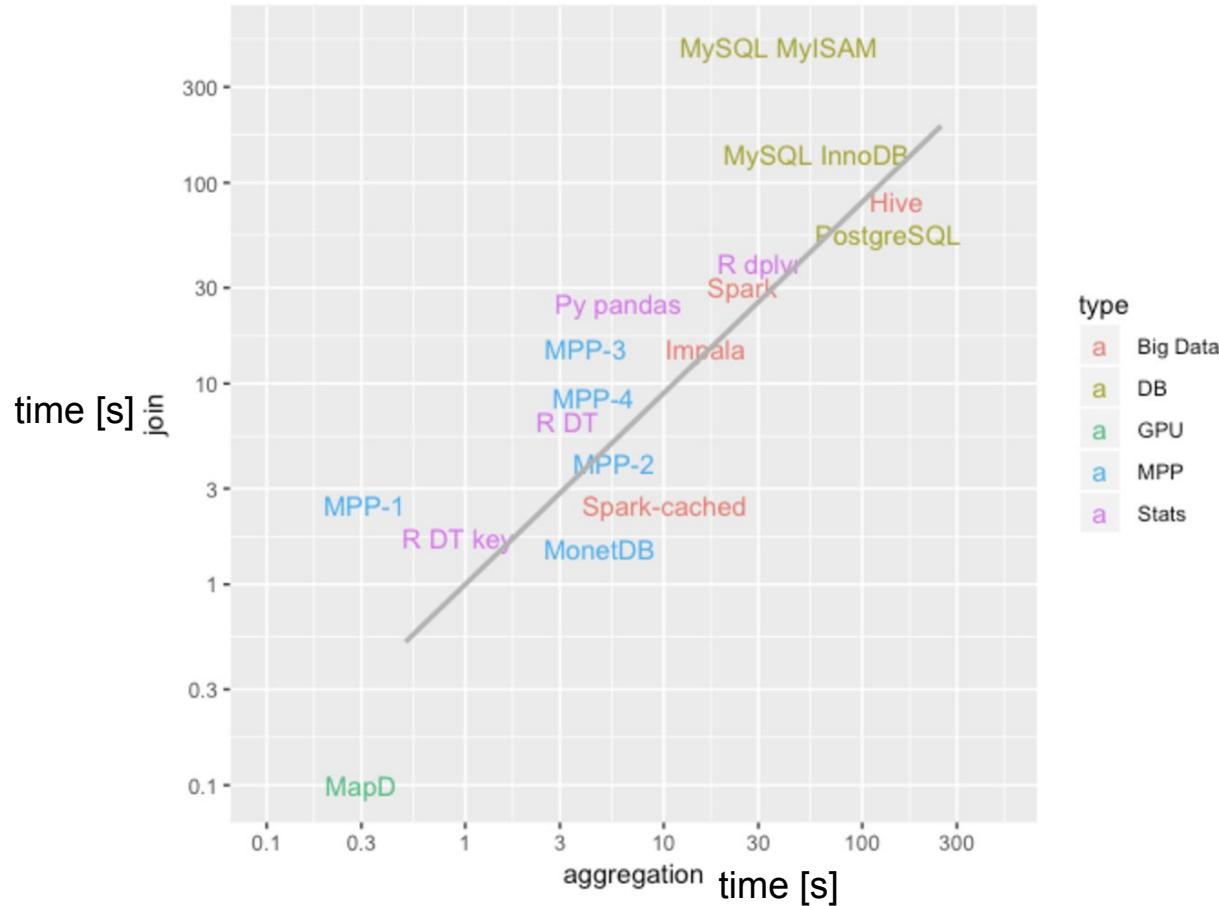
answered Nov 30 '09 at 15:48



JD Long

27.5k ⚡ 31 🔍 135 🏷 216







szilard / dataset-sizes-kdnuggets

quantile	value
50%	30 GB
60%	120 GB
70%	0.5 TB
80%	2 TB
90%	8 TB

(largest data analyzed)



szilard / dataset-sizes-kdnuggets

quantile	value
50%	30 GB
60%	120 GB
70%	0.5 TB
80%	2 TB
90%	8 TB

(largest data analyzed)

Michael Stonebraker | Big Data is (at least) Four Different Problems

Column Store Architecture (now)

E.g. Vertica, Sybase IQ, Microsoft PDW, DB2-BLU, RedShift, ...

DBaaS Database Group

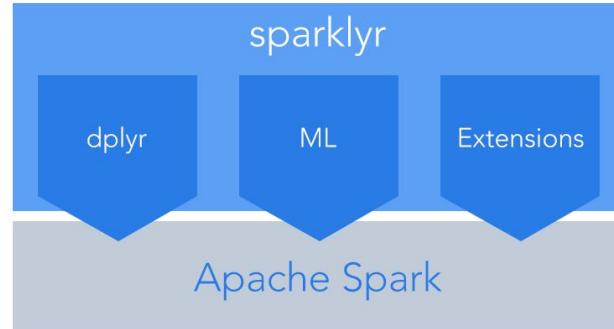
Computer Science and Artificial Intelligence Lab



szilard / dataset-sizes-kdnuggets

quantile	value
50%	30 GB
60%	120 GB
70%	0.5 TB
80%	2 TB
90%	8 TB

(largest data analyzed)



A video player interface showing a man in a red shirt speaking. The video is titled "Michael Stonebraker | Big Data is (at least) Four Different Problems". The video progress bar shows 3:13 / 55:52. The video is set against a background of a presentation slide titled "Column Store Architecture (now)" showing five green vertical bars labeled "COLUMN 1", "COLUMN 2", "COLUMN 3", "COLUMN 4", and "COLUMN K". Below the bars is the text "E.g. Vertica, Sybase IQ, Microsoft PDW, DB2-BLU, RedShift, ...". The DB logo and "Database Group" text are visible in the bottom right corner of the slide.

Michael Stonebraker | Big Data is (at least) Four Different Problems

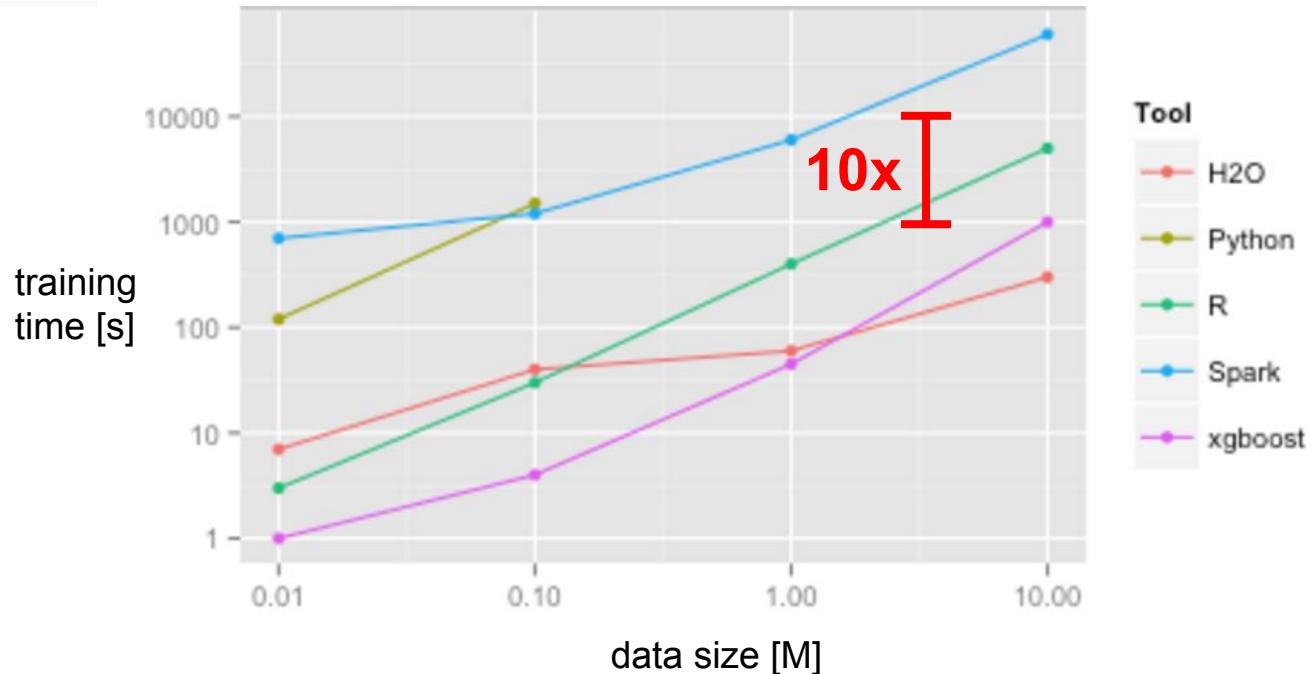


szilard / benchm-ml

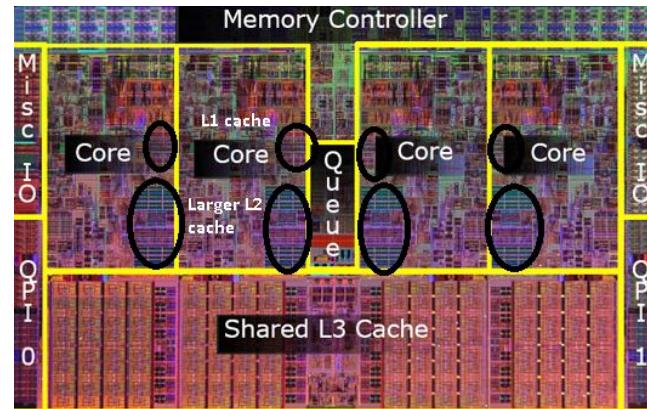
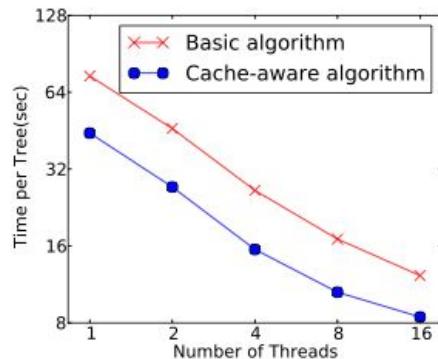


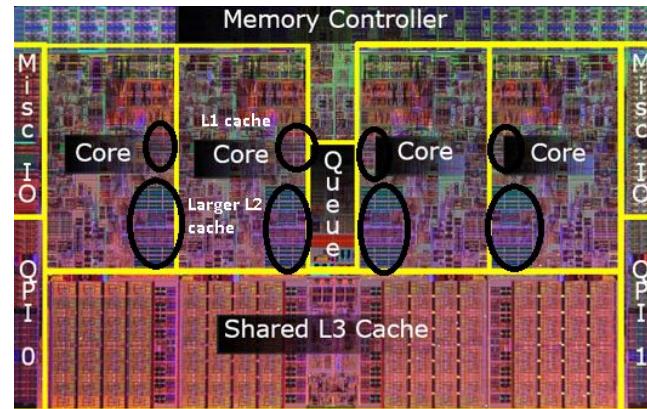
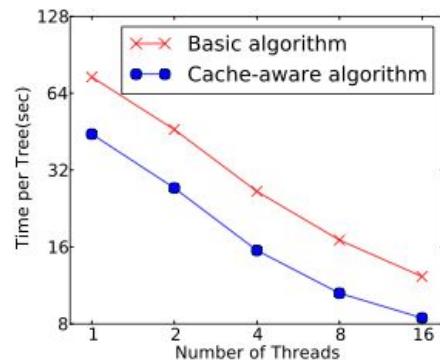
Gradient Boosting Machines

szilard / benchm-ml



XGBoost: A Scalable Tree Boosting System



XGBoost: A Scalable Tree Boosting System

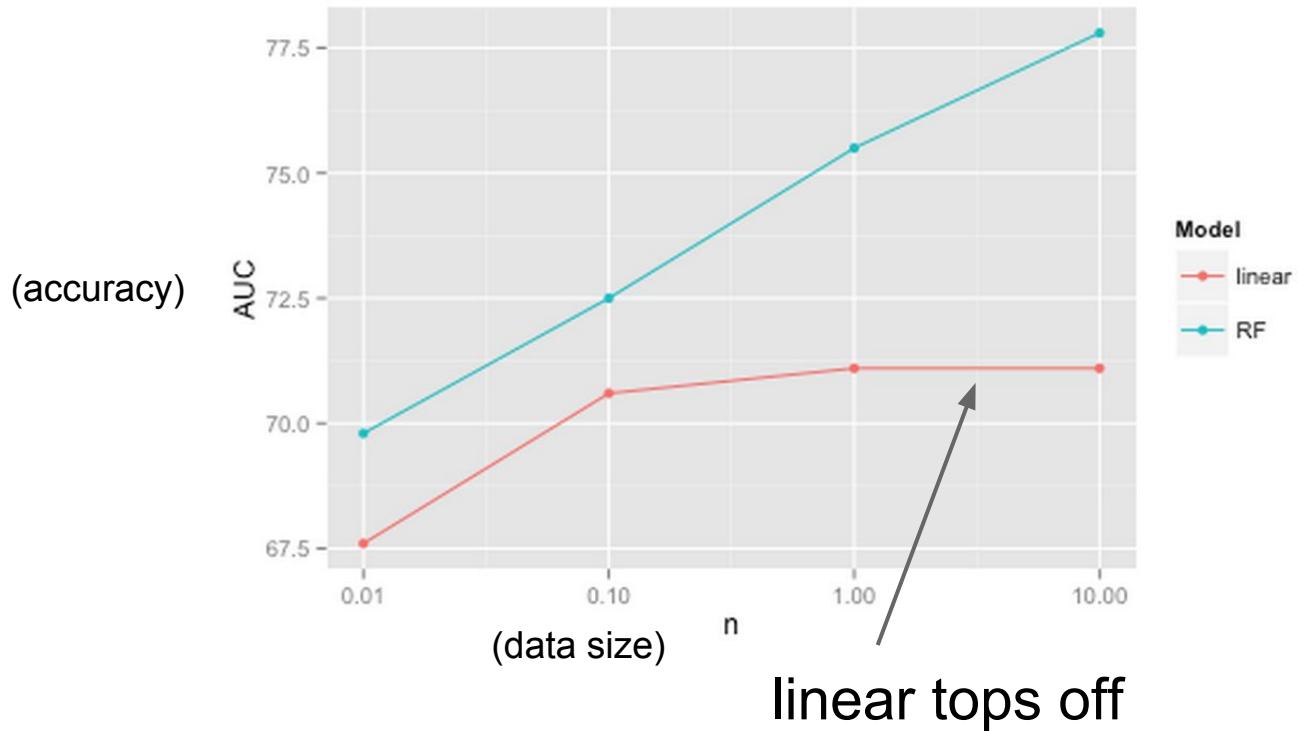
← → C <https://cran.r-project.org/web/pa>

xgboost: Extreme Gradient Boosting

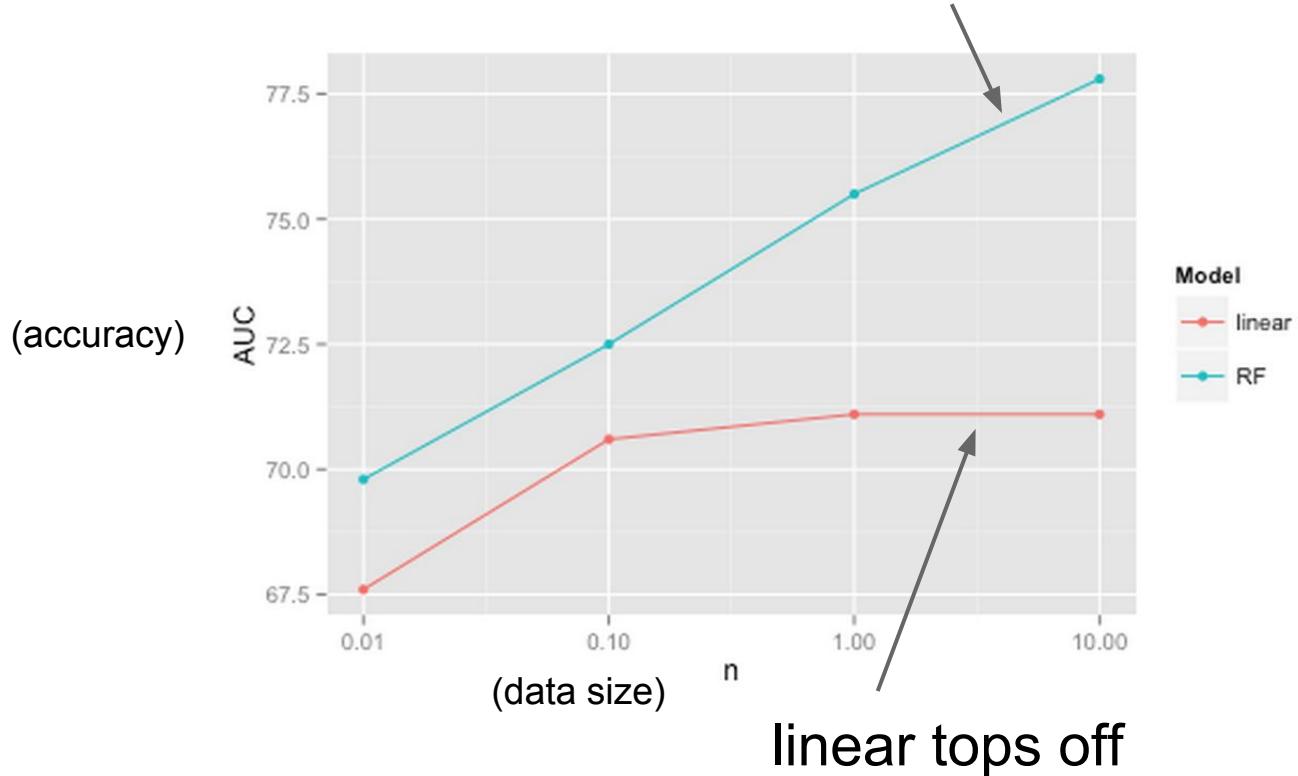
← → C <https://cran.r-project.org/>

h2o: R Interface for H2O



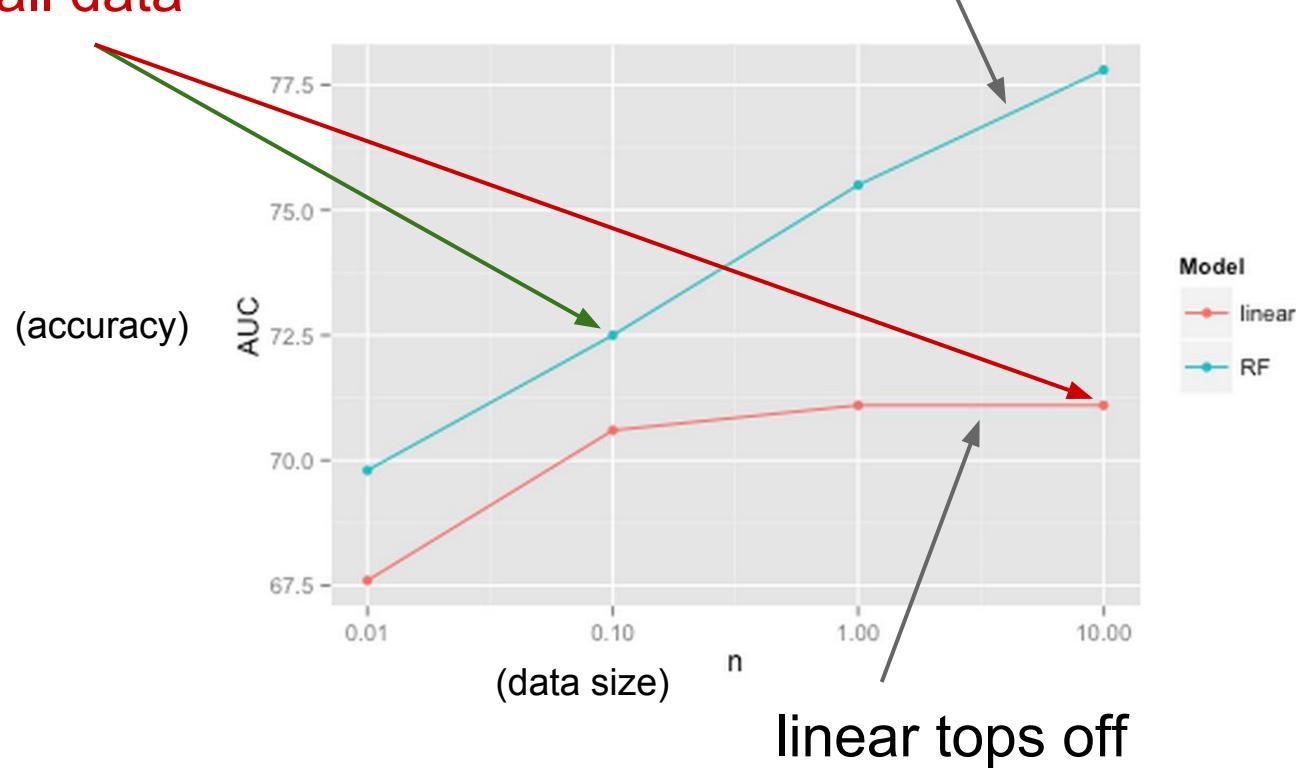


more data & better algo



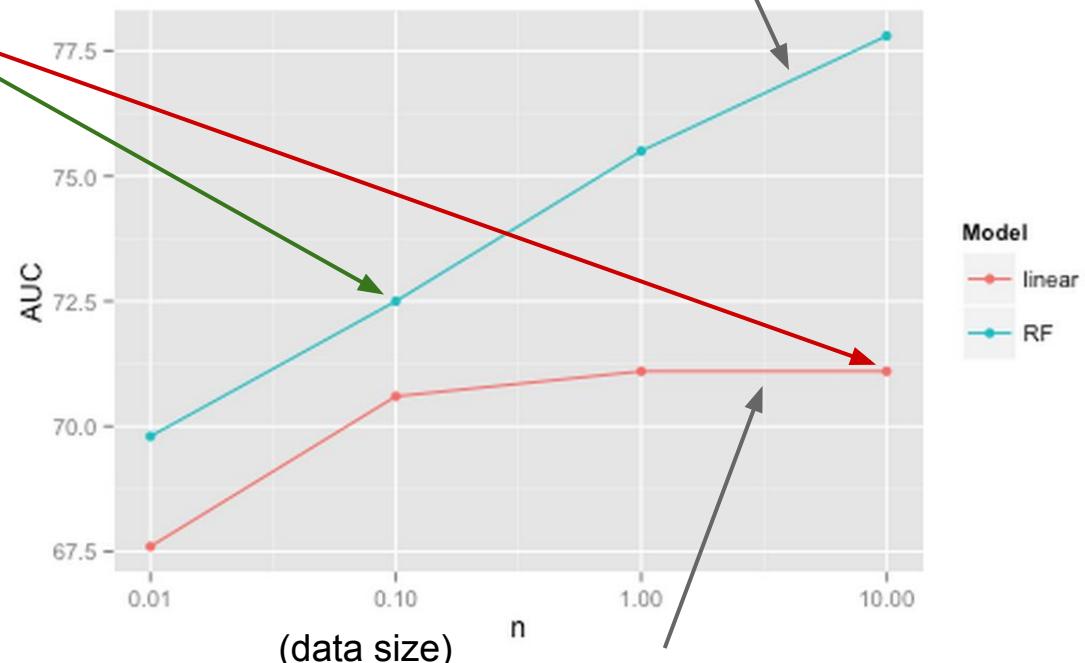
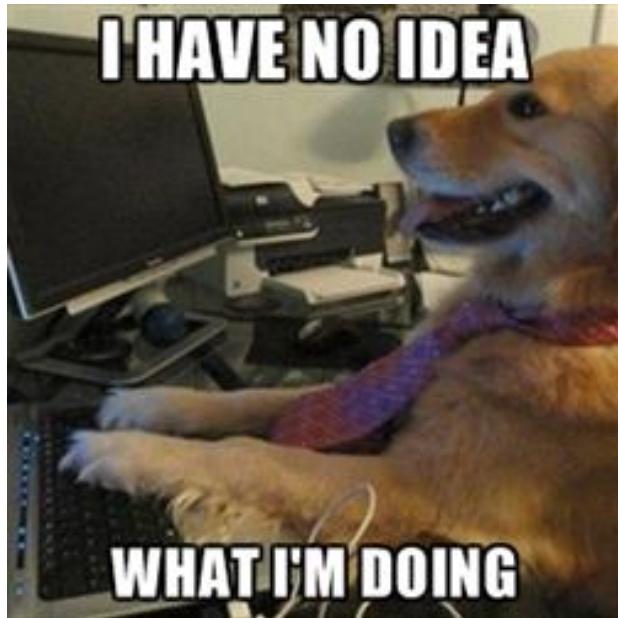
random forest on 1% of data
beats linear on all data

more data & better algo



random forest on 1% of data
beats linear on all data

more data & better algo



linear tops off





The Visual Display of Quantitative Information

EDWARD R. TUFTE

Trevor Hastie
Robert Tibshirani
Jerome Friedman

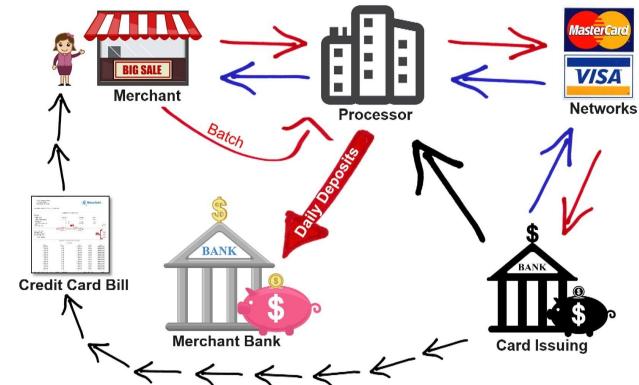
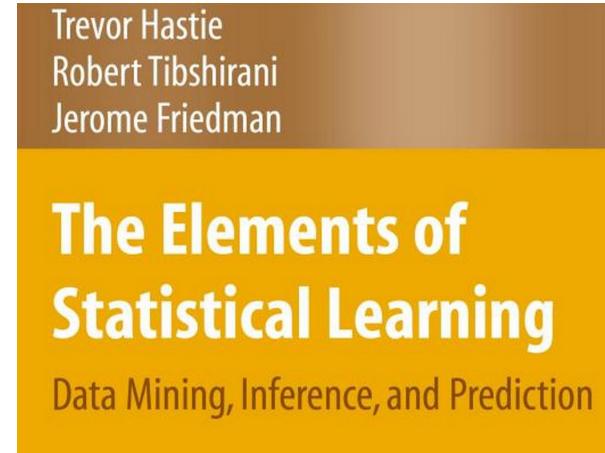
The Elements of Statistical Learning

Data Mining, Inference, and Prediction



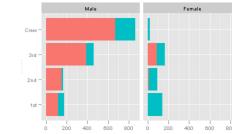
The Visual Display of Quantitative Information

EDWARD R. TUFTE





{ api }



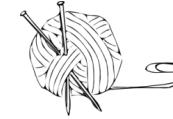
8,576
active packages

data.table

Seamless R
and C++
Integration
with Rcpp



R Studio®

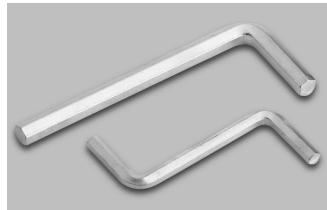
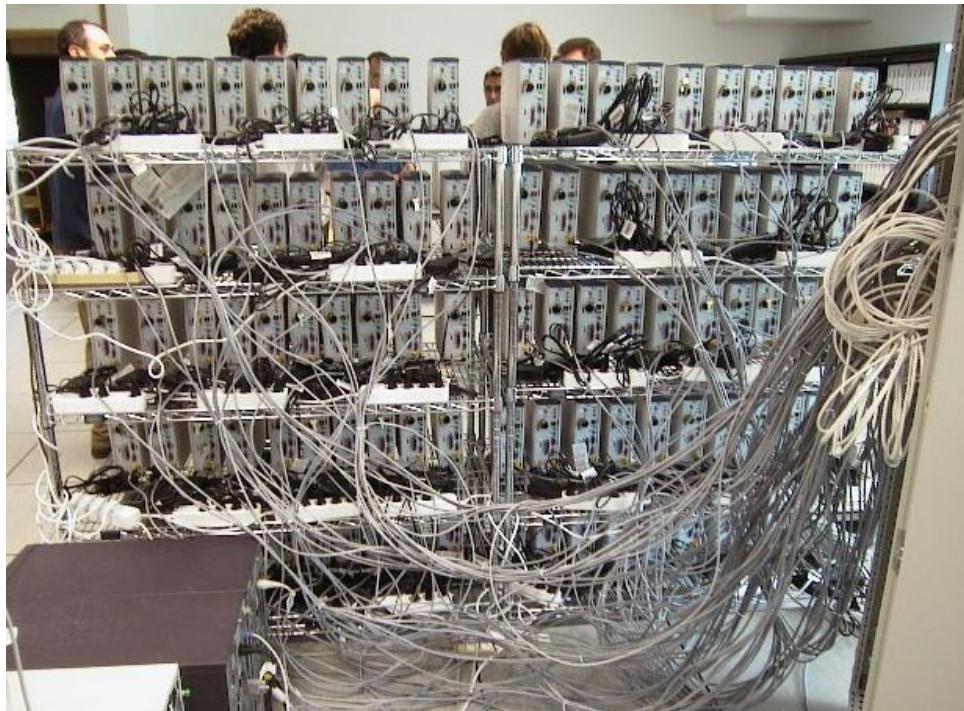


Meetup

stack overflow

use R!

R consortium

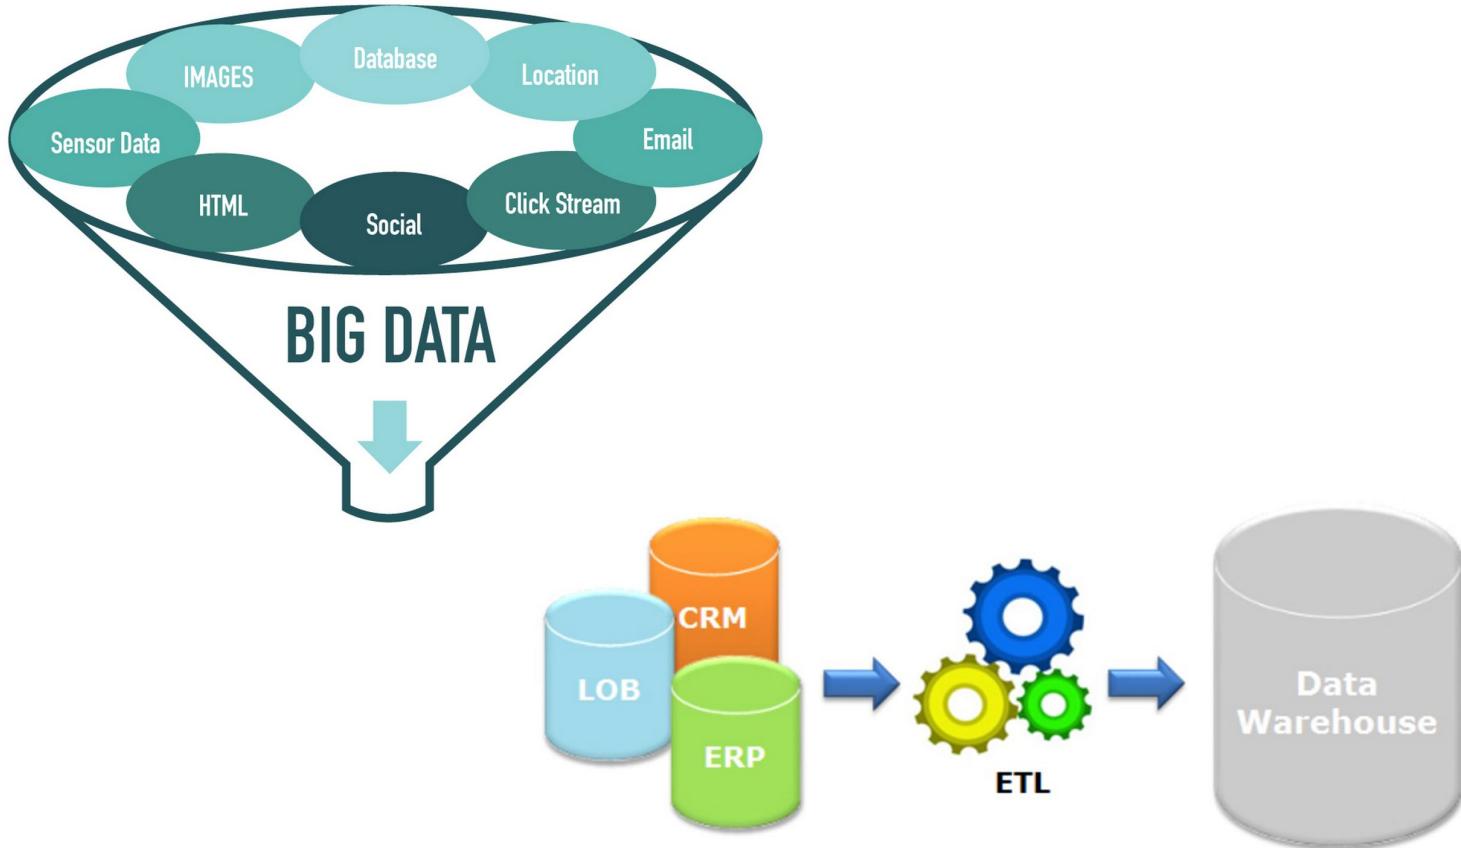


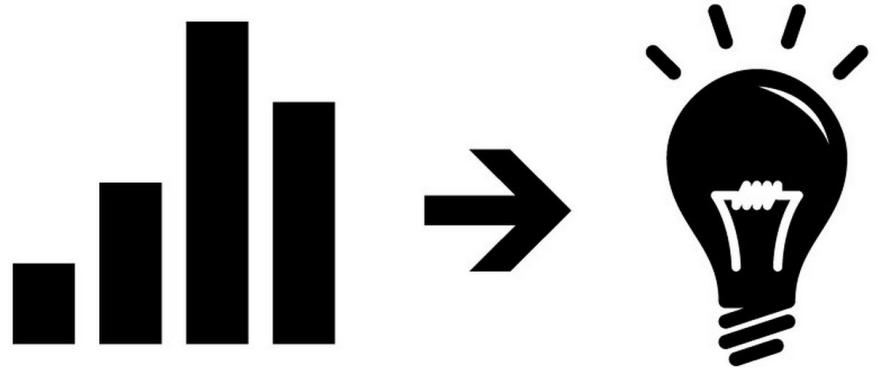
6	6	2	7	5	7	2	6	1				
4	2	3	1	4	4	4	7	7	5	3	2	1
0	4	0	3	3	9	0	5	9	7	8	3	9
8	0	SYSTEM FAILURE	4	1	0							
8	3	2	3	9	8	0	3	6	0	5	2	8
2	5	1	9	8	7	8	2	4	4	3	4	0
6	6	5	6	6	5	6	4	7	8	1	3	9

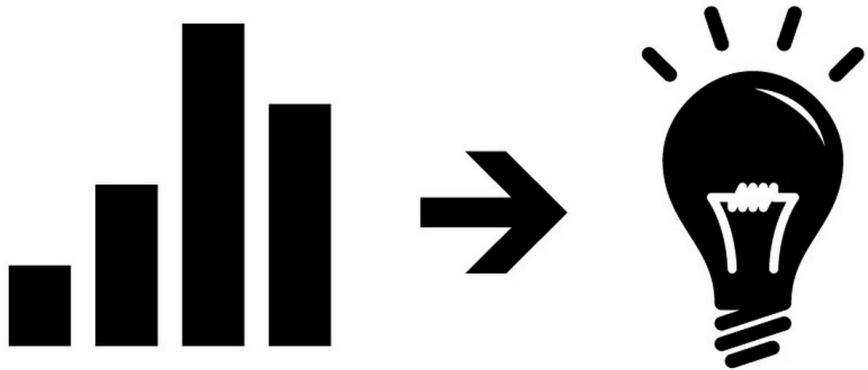
I DON'T ALWAYS ANALYZE
MY DATA, BUT WHEN I DO



I SPEND ALL MY TIME
FIGHTING THE TOOLS







BE PRODUCTIVE.

Summary / Tips for analyzing “big” data:

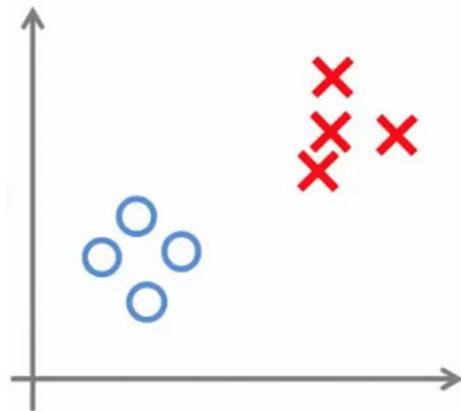
- Get lots of RAM (physical/ cloud)
- Use R/Python and high performance packages (e.g. data.table, xgboost)
- Do data reduction in database (analytical db/ big data system)
- (Only) distribute embarrassingly parallel tasks (e.g. hyperparameter search for machine learning)
- Let engineers (store and) ETL the data (“scalable”)
- Use statistics/ domain knowledge/ thinking
- Use “big data tools” only if the above tips not enough

Example #2

?



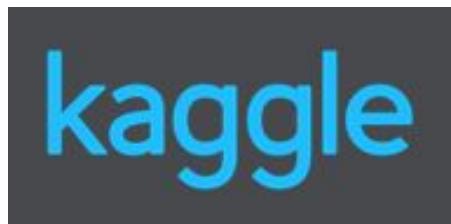
open source



(10,000,000 rows)

I usually use other people's code [...] I can find open source code for what I want to do, and my time is much better spent doing research and feature engineering -- Owen Zhang

<http://blog.kaggle.com/2015/06/22/profiling-top-kagglers-owen-zhang-currently-1-in-the-world/>



MASTER ?

1st
/328,471

176,181.4 points
Joined 4 years ago
†Ranking method changed 13 May 2015 (?)

A portrait photograph of Owen Zhang, a man with glasses and a smile, wearing a light-colored shirt.



[szilard / benchm-ml](#)



binary classification, 10M records
numeric & categorical features, non-sparse

MODEL	1ST	2ND	AVG	1ST	2ND
BST-DT	0.580	0.228	RF	0.727	0.207
RF	0.390	0.525	ANN	0.053	0.172
BAG-DT	0.030	0.232	BSTD T	0.059	0.228
SVM	0.000	0.008	SVM	0.043	0.195
ANN	0.000	0.007	LR	0.089	0.132
KNN	0.000	0.000	BAGDT	0.002	0.012
BST-STMP	0.000	0.000	KNN	0.023	0.045
DT	0.000	0.000	BSTST	0.004	0.009
LOGREG	0.000	0.000	PRC	0	0
NB	0.000	0.000	NB	0	0

An Empirical Comparison of Supervised Learning Algorithms

<http://www.cs.cornell.edu/~alexn/papers/empirical.icml06.pdf>

An Empirical Evaluation of Supervised Learning in High Dimensions

<http://lowrank.net/nikos/pubs/empirical.pdf>

MODEL	1ST	2ND
BST-DT	0.580	0.228
RF	0.390	0.525
BAG-DT	0.030	0.232
SVM	0.000	0.008
ANN	0.000	0.007
KNN	0.000	0.000
BST-STMP	0.000	0.000
DT	0.000	0.000
LOGREG	0.000	0.000
NB	0.000	0.000

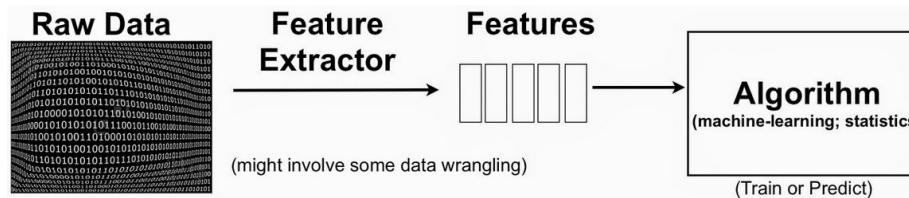
AVG	1ST	2ND
RF	0.727	0.207
ANN	0.053	0.172
BSTD T	0.059	0.228
SVM	0.043	0.195
LR	0.089	0.132
BAGDT	0.002	0.012
KNN	0.023	0.045
BSTST	0.004	0.009
PRC	0	0
NB	0	0

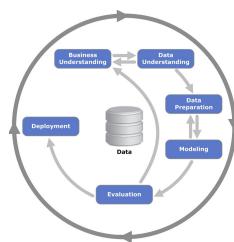
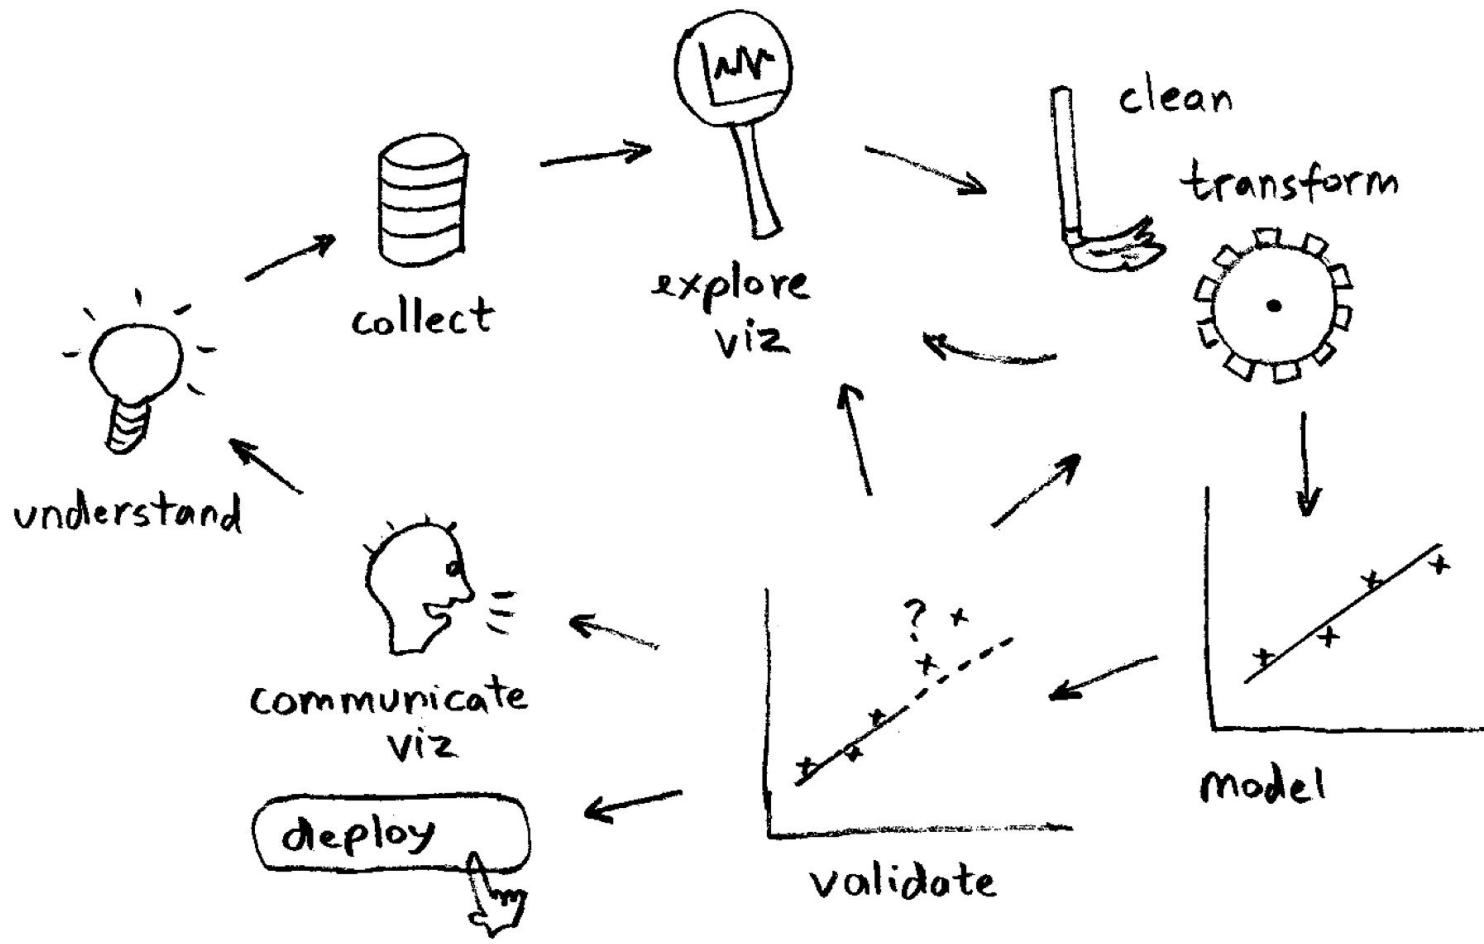
An Empirical Comparison of Supervised Learning Algorithms

<http://www.cs.cornell.edu/~alexn/papers/empirical.icml06.pdf>

An Empirical Evaluation of Supervised Learning in High Dimensions

<http://lowrank.net/nikos/pubs/empirical.pdf>







open source

- R packages
- Python scikit-learn
- Vowpal Wabbit
- H2O
- xgboost
- Spark MLlib
- a few others



open source

- R packages 30%
- Python scikit-learn 40%
- Vowpal Wabbit 8%
- H2O 10%
- xgboost 8%
- Spark MLlib 6%
- a few others



open source

- R packages **30%** *Lightning-fast*
- Python scikit-learn **40%**
- Vowpal Wabbit **8%** *Big Data*
- H2O **10%**
- xgboost **8%** *Large-scale*
- Spark MLlib **6%**
- a few others *optimized*

Distributed

Yes, it is possible to apply **big data!**

parallelized

Scalable and fast

scalable performant machine learning

built for scale

Destination	Gate	Time	On Time
PARIS	A43	12:00	ON TIME
FRANKFURT	A15	12:10	ON TIME
NEW YORK	B08	12:25	ON TIME
BRUSSELS	A21	12:30	ON TIME
ROME	A30	12:30	ON TIME
BOSTON	B01	12:35	ON TIME
LONDON	A19	12:40	ON TIME
RIO DE JANEIRO	B13	12:45	ON TIME
MADRID	A26	12:45	ON TIME
ATHENS	A37	12:50	ON TIME
STOCKHOLM	A40	13:00	ON TIME
DUBLIN			

Destination	Gate	Time	ON	TIME
PARIS	A43	12:00	ON	TIME
FRANKFURT	A15	12:10	ON	TIME
NEW YORK	B08	12:25	ON	TIME
BRUSSELS	A21	12:30	ON	TIME
ROME	A30	12:30	ON	TIME
BOSTON	B01	12:35	ON	TIME
LONDON	A19	12:40	ON	TIME
RIO DE JANEIRO	B13	12:45	ON	TIME
MADRID	A26	12:45	ON	TIME
ATHENS	A37	12:50	ON	TIME
STOCKHOLM	A40	13:00	ON	TIME
DUBLIN				

EC2



$n = 10K, 100K, 1M, 10M, 100M$

Training time

RAM usage

AUC

CPU % by core

read data, pre-process, score test data

$n = 10K, 100K, 1M, 10M, 100M$

Training time

RAM usage

AUC

CPU % by core

read data, pre-process, score test data





szilard / **benchm-ml**

branch: **master** ▾

benchm-ml / **2-rf** / +

xgboost improve



szilard authored 19 days ago

..

1.R

2.py

4-h2o-v3.R

4-h2o.R







[szilard / benchm-ml](#)

★ Star

1,117

Simple/limited/incomplete benchmark



[szilard / benchm-ml](#)



1,117

Simple/limited/incomplete benchmark

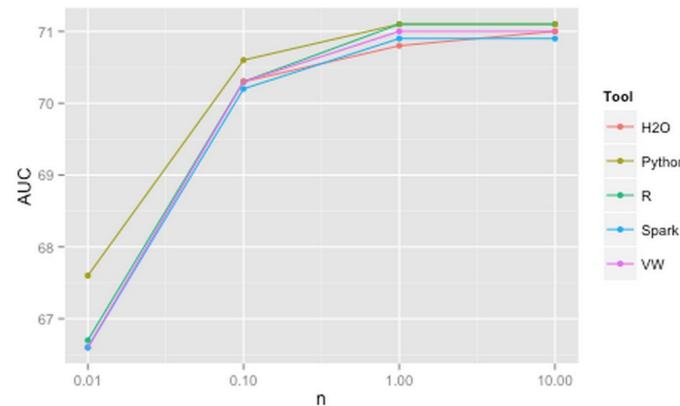
All benchmarks are wrong, but some are useful

LINEAR



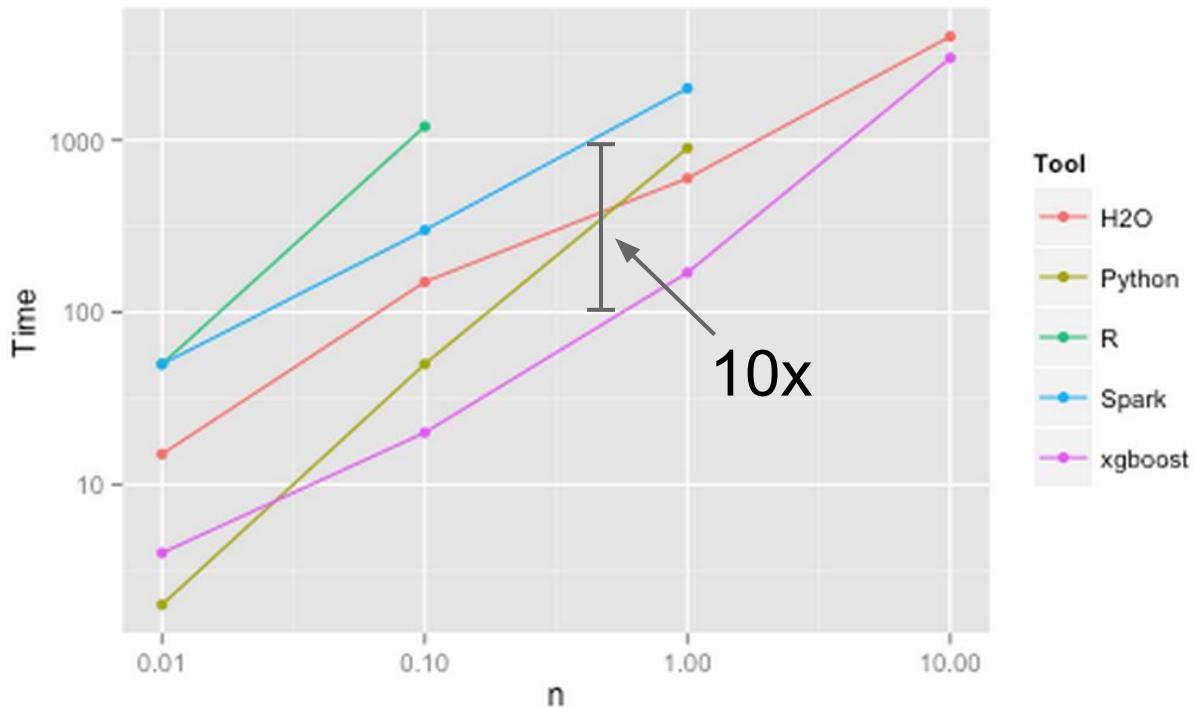


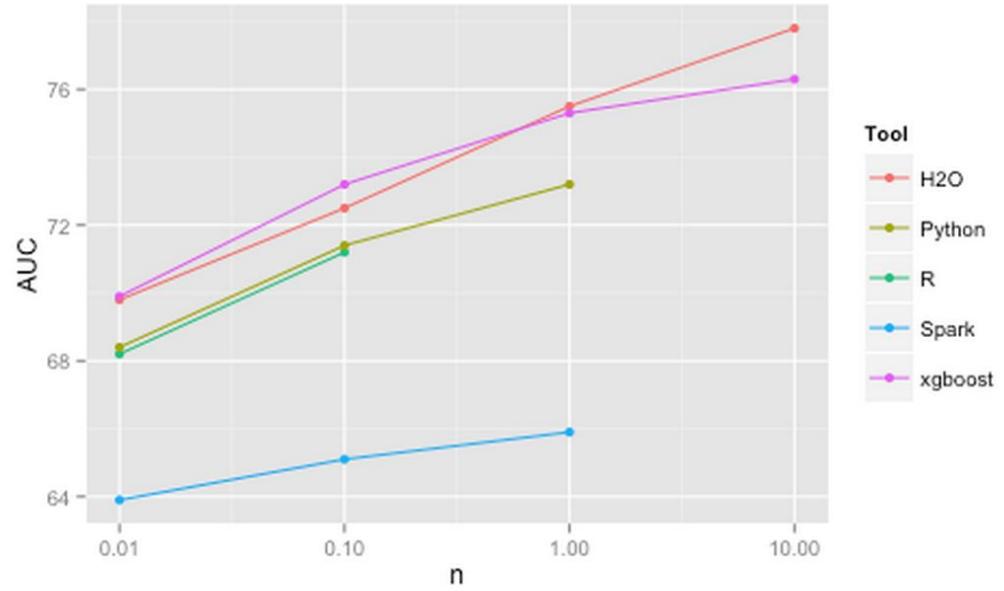
Tool	n	Time (sec)	RAM (GB)
R	10K	0.1	1
	100K	0.5	1
	1M	5	1
	10M	90	5



The main conclusion here is that it is trivial to train linear models even for $n = 10M$ rows virtually in any of these tools on a single machine in a matter of seconds. H2O and VW are the most memory efficient

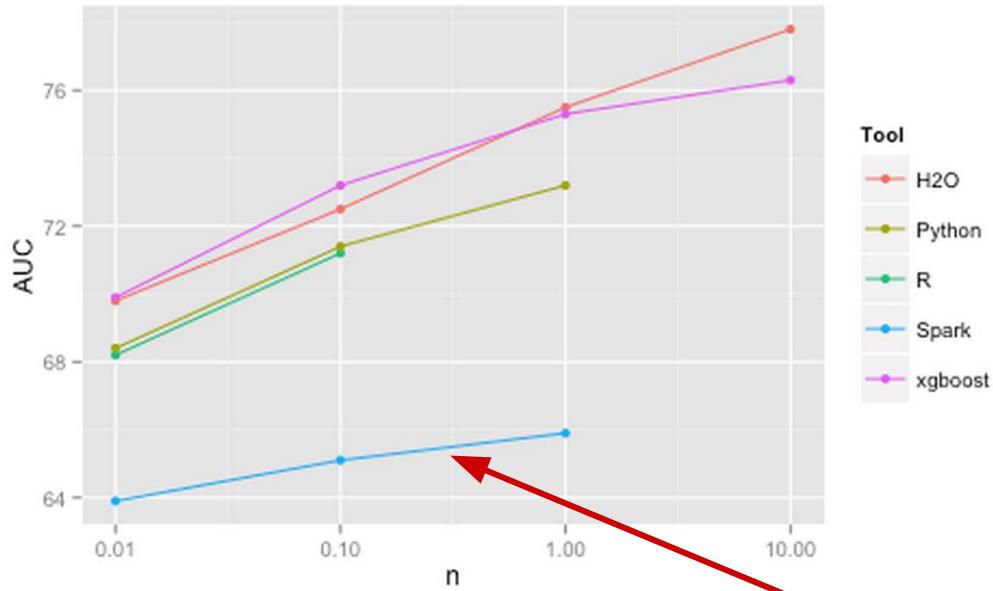






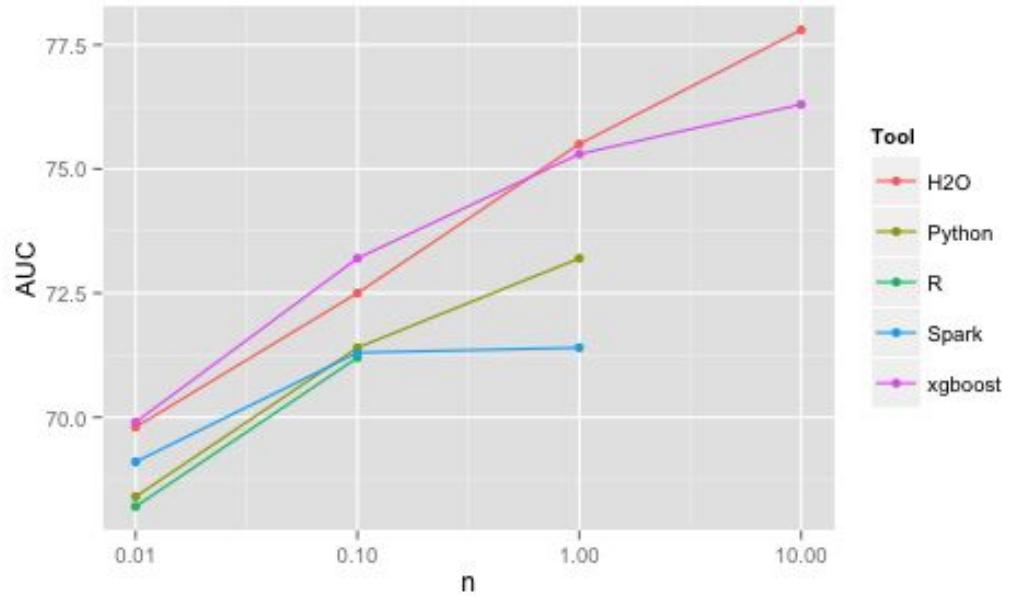
Tool

- H2O
- Python
- R
- Spark
- xgboost



Joseph Bradley

AUC/accuracy: The AUC issue appears to be caused by MLlib tree ensembles aggregating votes, rather than class probabilities, as you suggested. I re-ran your test using class probabilities (which can be aggregated by hand), and then got the same AUC as other libraries. We're planning on including this fix in Spark 1.5 (and thanks for providing some evidence of its importance!).



Tool

- H2O
- Python
- R
- Spark
- xgboost

Timing: I didn't try to reproduce your results yet, but have a few thoughts. The main issue with MLlib's tree implementation is that it is optimized for training shallow trees, following the PLANET project. We're working on an alternative implementation geared towards training deep trees, hopefully

aimed at Spark 1.5 or 1.6. One big benefit of running on top of Spark is that there is constant work on improving the underlying system, which MLlib will benefit from. In particular, the JVM memory management issues will improve as project Tungsten (which can be Googled) progresses.

<http://datascience.la/benchmarking-random-forest-implementations/#comment-53599>



[Download](#) [Libraries](#) ▾ [Documentation](#)

Spark Release 2.0.0



[Download](#) [Libraries](#) ▾ [Documentation](#)

Spark Release 2.0.0

DataFrame-based API is primary API



[Download](#) [Libraries](#) ▾ [Documentation](#)

Spark Release 2.0.0

DataFrame-based API is primary API

Spark random forest issues #19

[① Open](#) szilard opened this issue on Jul 23, 2015 · 15 comments



szilard commented 4 days ago

It is slower than before

MLlib 1.5 – 250 sec
ML 2.0 – 400 sec



[Download](#) [Libraries](#) ▾ [Documentation](#)

Spark Release 2.0.0

DataFrame-based API is primary API

Spark random forest issues #19

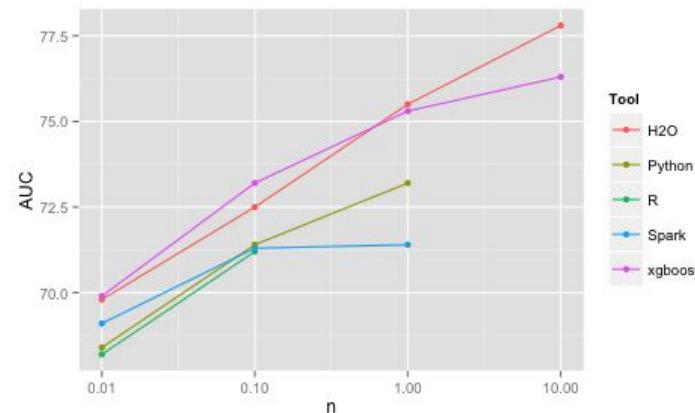
[Open](#) szilard opened this issue on Jul 23, 2015 · 15 comments

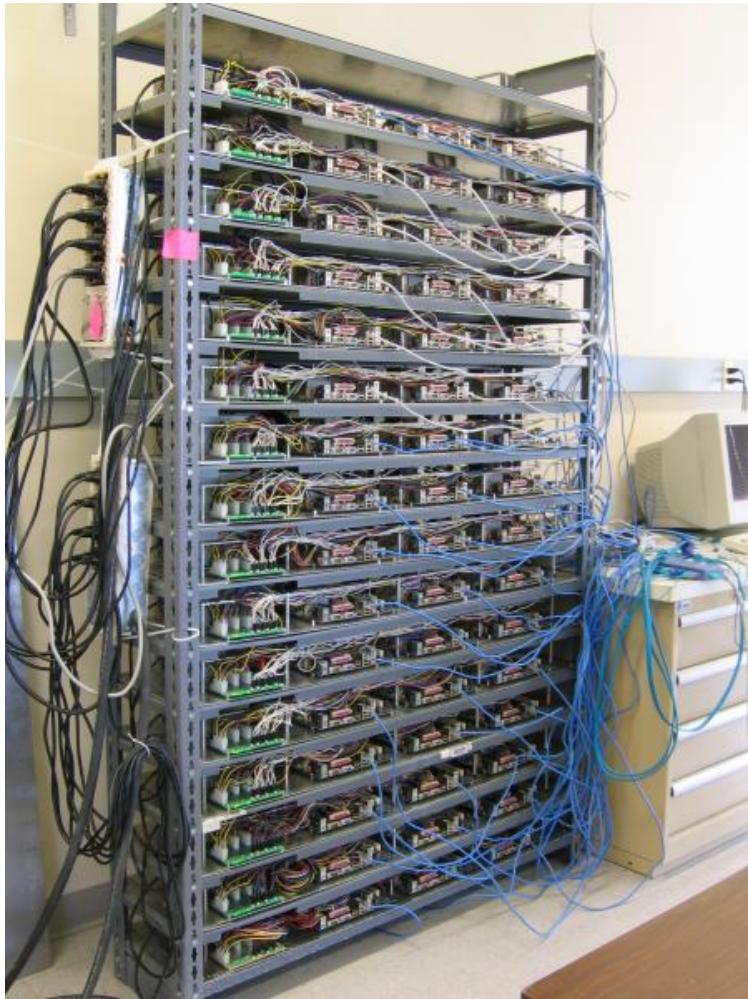


szilard commented 4 days ago

It is slower than before

MLlib 1.5 – 250 sec
ML 2.0 – 400 sec



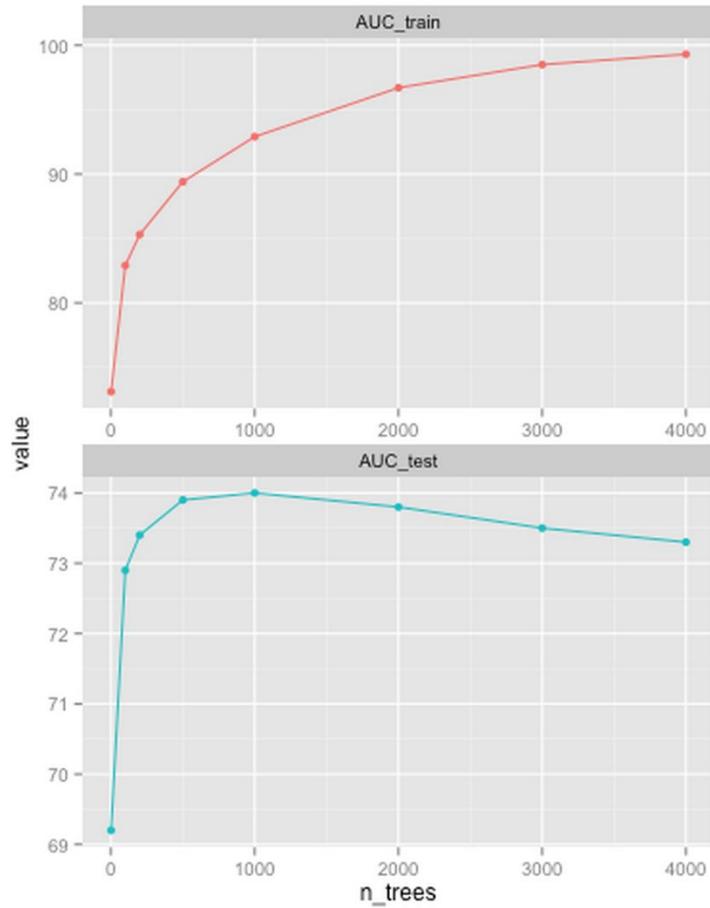


ntree	depth	nbins	mtries	Time (hrs)	AUC
500	20	20	-1 (2)	1.2	77.8
500	50	200	-1 (2)	4.5	78.9
500	50	200	3	5.5	78.9
5000	50	200	-1 (2)	45	79.0
500	100	1000	-1 (2)	8.3	80.1

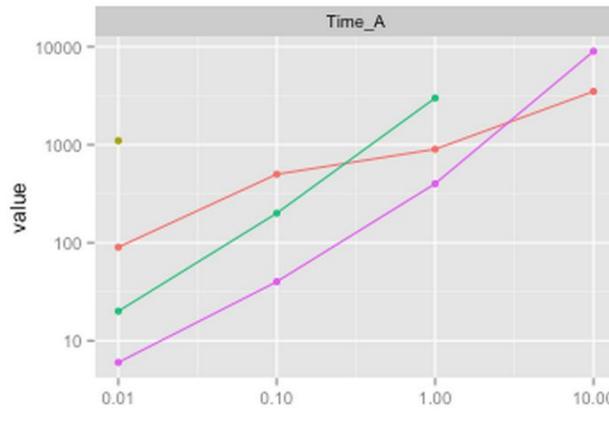
ntree	depth	nbins	mtries	Time (hrs)	AUC
500	20	20	-1 (2)	1.2	77.8
500	50	200	-1 (2)	4.5	78.9
500	50	200	3	5.5	78.9
5000	50	200	-1 (2)	45	79.0
500	100	1000	-1 (2)	8.3	80.1

Best linear: 71.1

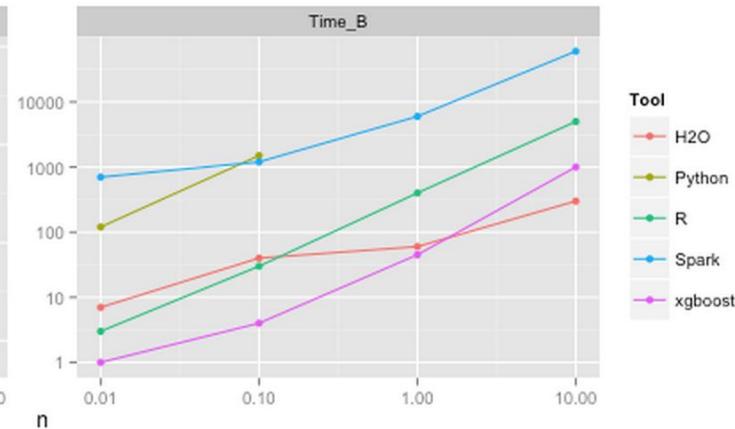




learn_rate = 0.01, max_depth = 16, n_trees = 1000

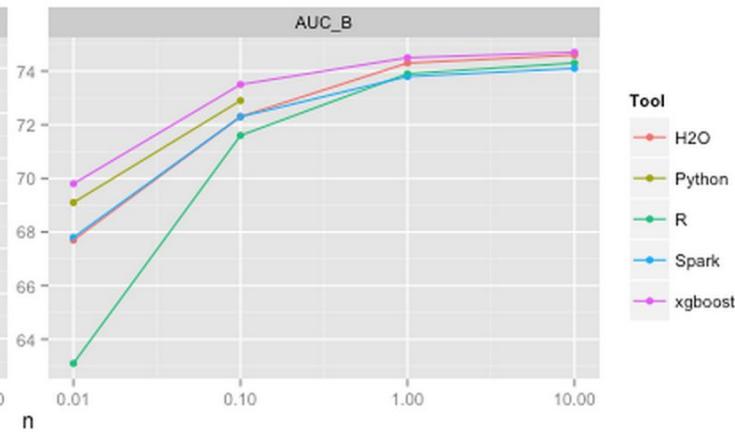
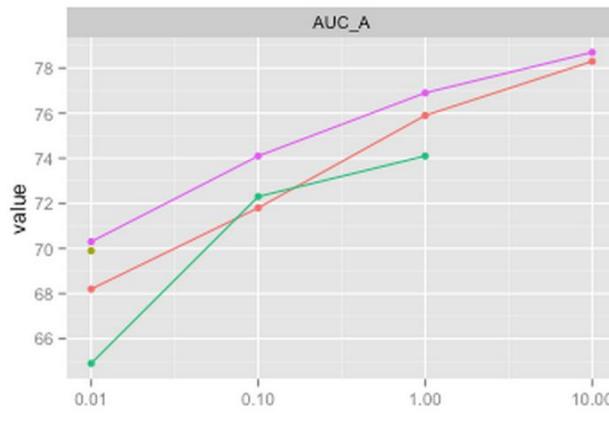


learn_rate = 0.1, max_depth = 6, n_trees = 300



Tool

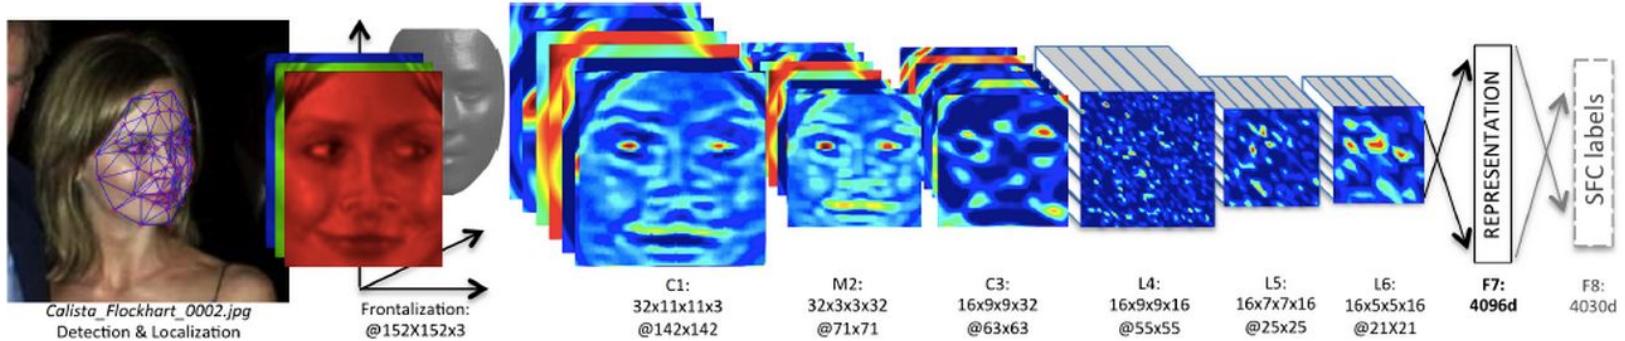
- H2O
- Python
- R
- Spark
- xgboost

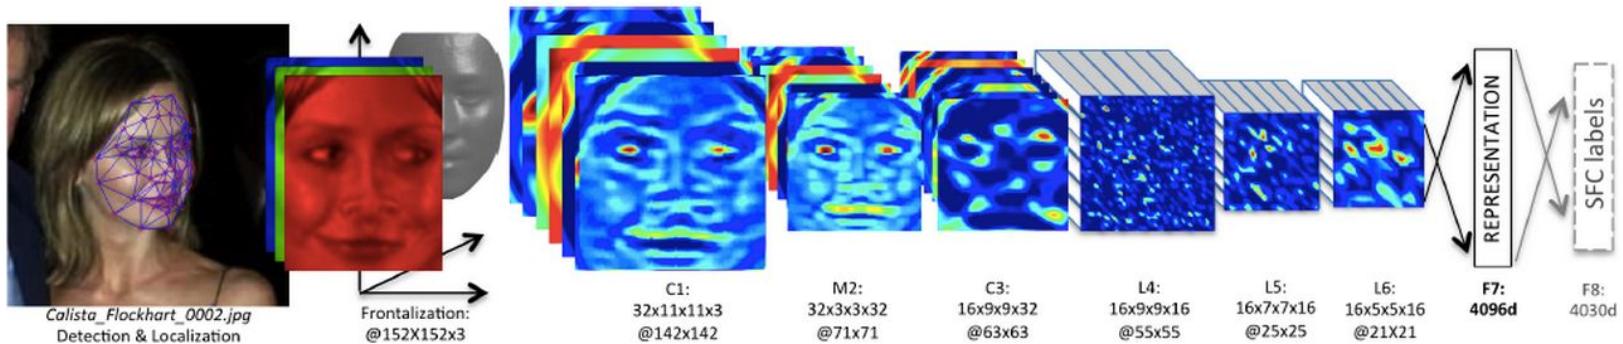


Tool

- H2O
- Python
- R
- Spark
- xgboost

Tool	Time (hr)	AUC
H2O	7.5	79.8
H2O-3	9.5	81.2
xgboost	14	81.1





Params	AUC	Time (s)	Epochs
default: activation = "Rectifier", hidden = c(200,200)	73.1	270	1.8
hidden = c(50,50,50,50), input_dropout_ratio = 0.2	73.2	140	2.7

...

ADADELTA rho = 0.95, epsilon = 1e-06	71.1	240	1.7
rho = 0.999, epsilon = 1e-08	73.3	270	1.9
adaptive = FALSE default: rate = 0.005, decay = 1, momentum = 0	73.0	340	1.1

MY DATA



IS BIGGER THAN YOURS

Larger Data Sizes (on a Single Server)



Big Data Borat @BigDataBorat · 22 Nov 2013

cat data data data data data data > bigdata

Larger Data Sizes (on a Single Server)



Big Data Borat @BigDataBorat · 22 Nov 2013

cat data data data data data data > bigdata

Linear models, 100M rows:

Tool	Time[s]	RAM[GB]
R	1000	60
Spark	160	120
H2O	40	20
VW	150	

Larger Data Sizes (on a Single Server)



Big Data Borat @BigDataBorat · 22 Nov 2013

cat data data data data data data > bigdata

Linear models, 100M rows:

Tool	Time[s]	RAM[GB]
R	1000	60
Spark	160	120
H2O	40	20
VW	150	

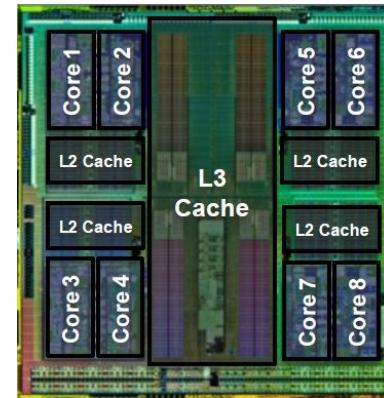
Linear models, 1B rows:

Tool	Time[s]	RAM[GB]
H2O	500	100
VW	1400	

Larger Data Sizes (on a Single Server)

RF/GBM, 100M rows:

Algo	Tool	Time[s]	Time[hr]	RAM[GB]
RF	H2O	40000	11	80
	xgboost	36000	10	60
GBM	H2O	35000	10	100
	xgboost	110000	30	50



Distributed Systems

H2O logistic runtime (sec):

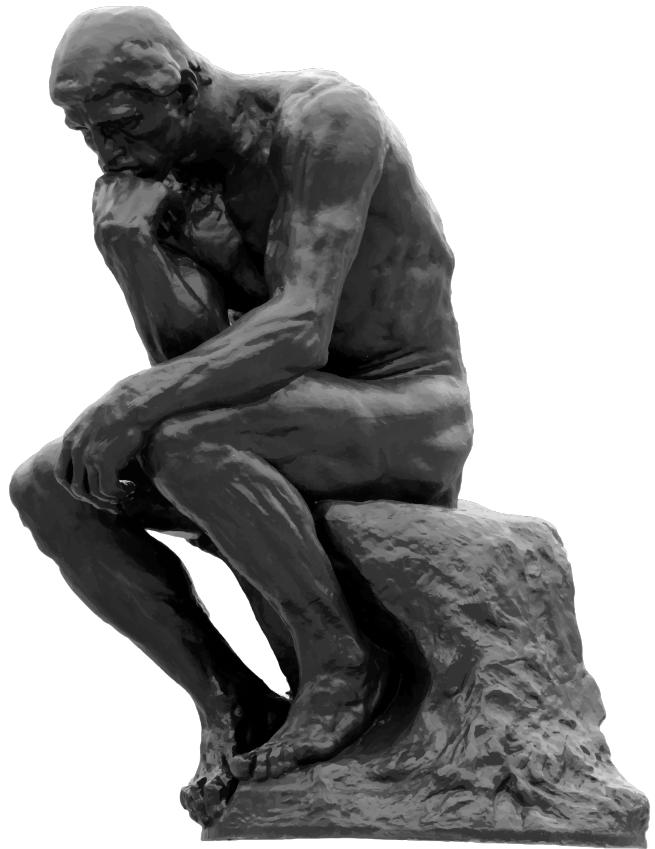
	1 node	5 nodes
100M	42	9.9
1B	480	101

H2O RF runtime (sec) (5 trees):

	1 node	5 nodes
10M	42	41
100M	405	122

Summary





Business Optimization



Business Optimization





✉️ spafka@gmail.com



🐦 [@DataScienceLA](https://twitter.com/DataScienceLA)

linkedin linkedin.com/in/szilard



github github.com/szilard